

# Task 2: Deep Learning based Quark-Gluon Classification

---

● **Data Preparation:** Please train your model on 80% of the data and evaluate on the remaining 20%. Please make sure not to overfit on the test dataset - it will be checked with an independent sample.

● **Model Training:** Train a **VGG13** model and another model of your choice.

---

```
In [1]: import os
import numpy as np
import matplotlib.pyplot as plt
```

```
In [2]: import gc
import pyarrow.parquet as pq
from random import shuffle
```

```
In [3]: import torch
from torch import nn
from torch import optim

from torchvision import transforms as T
from torch.utils.data import Dataset, DataLoader, TensorDataset, random_split

from torchvision import models
```

```
/home/harkhymadhe/miniforge3/lib/python3.11/site-packages/torchvision/io/image.py:13: UserWarning: Failed to load image Python extension: '/home/harkhymadhe/miniforge3/lib/python3.11/site-packages/torchvision/image.so: undefined symbol: _ZN3c106detail23torchInternalAssertFailEPKcS2_jS2_RKSs' If you don't plan on using image functionality from `torchvision.io`, you can ignore this warning. Otherwise, there might be something wrong with your environment. Did you have `libjpeg` or `libpng` installed before building `torchvision` from source?
warn(
```

---

## Phase I: Data Preparation

**Aim:** Please train your model on 80% of the data and evaluate on the remaining 20%. Please make sure not to overfit on the test dataset - it will be checked with an independent sample.

---

First, the parquet files are downloaded and stored in the **./dataset/** folder.

```
In [4]: # File paths
file1 = "dataset/QCDToGGQQ_IMGjet_RH1all_jet0_run0_n36272.test.snappy.parquet"
file2 = "dataset/QCDToGGQQ_IMGjet_RH1all_jet0_run1_n47540.test.snappy.parquet"
file3 = "dataset/QCDToGGQQ_IMGjet_RH1all_jet0_run2_n55494.test.snappy.parquet"
```

A little bit of experimentation showed that loading the Parquet files via **Pandas** or basic **PyArrow** was very inefficient and resulted in OOM errors. I attempt to bypass this via the more specialized **parquet** subpackage in **PyArrow**.

```
In [5]: # Load data files
class ParquetDataset(Dataset):
    def __init__(self, filename):
        self.parquet = pq.ParquetFile(filename)
        self.cols = None

    def __getitem__(self, index):
        data = self.parquet.read_row_group(index, columns=self.cols).to_pydict()
        data['X_jets'] = 1.*np.float32(data['X_jets'][0])#/data['mGG']
        data['X_jets'] = data['X_jets'][0][:80000]

        data = dict(data)
        return torch.as_tensor(np.expand_dims(data["X_jets"], axis = 0)), index

    def __len__(self):
        return self.parquet.num_row_groups

    @classmethod
    def from_files(cls, filenames):
        return ConcatDataset([cls(fname) for fname in filenames])
```

```
In [ ]:
```

Loading the Parquet data using the **ParquetDataset** class defined above is quite OK, but it makes actual file loading for multiple data points more cumbersome. A more efficient **BatchedParquetDataset** is implemented below:

```
In [6]: # Load data files
class BatchedParquetDataset(Dataset):
    def __init__(self, filename, batch_size):
        super().__init__()

        self.batch_size = batch_size
        self.parquet = pq.ParquetFile(filename)
        self.cols = None

        self.size = self.parquet.num_row_groups

        self.remainder = self.size % self.batch_size
```

```

        self.batch_indices = list(range(0, self.size, self.batch_size))
        self.batch_indices = list(
            zip(
                self.batch_indices,
                self.batch_indices[1:] + [self.batch_indices[-1] + (self.rem
            )
        )

    def __getitem__(self, index):
        indexes = range(*self.batch_indices[index])
        data = self.parquet.read_row_groups(indexes, columns=self.cols).columns

        image = torch.as_tensor(data[0].to_pylist())
        targets = torch.as_tensor(data[-1].to_pylist(), dtype = torch.long)

        return image, targets

    def __len__(self):
        return len(self.batch_indices)

    @classmethod
    def from_files(cls, filenames, batch_size):
        return ConcatDataset([cls(filename = fname, batch_size = batch_size)

```

```
In [7]: batch_size = 64
```

```
In [8]: batched_data = BatchedParquetDataset.from_files(batch_size = batch_size, fil
```

```
In [9]: sample = batched_data[0]
```

```
In [10]: sample[0].shape
```

```
Out[10]: torch.Size([64, 3, 125, 125])
```

```
In [11]: sample[1].shape
```

```
Out[11]: torch.Size([64])
```

```
In [12]: len(batched_data) * batch_size
```

```
Out[12]: 139392
```

```
In [ ]:
```

```
In [13]: train_data, test_data = random_split(batched_data, lengths = [.8, .2])
```

```

class BatchedDataLoader(DataLoader):
    def __init__(self, args, **kwargs):
        super().__init__(args, **kwargs)

```

```

    def __iter__(self):
        return iter(super())

```

```
def __next__(self):  
    return
```

```
In [14]: train_dl = DataLoader(  
        dataset = train_data,  
        batch_size = 1,  
        shuffle = True,  
        num_workers = 4,  
        pin_memory = True  
    )  
  
    test_dl = DataLoader(  
        dataset = test_data,  
        batch_size = 1,  
        shuffle = True,  
        num_workers = 4,  
        pin_memory = True  
    )
```

```
In [ ]:
```

```
In [15]: # Set device  
DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```
In [16]: gc.collect()
```

```
Out[16]: 20
```

---

## Phase II: Model Training

**Aim:** Train a VGG13 model and another model of your choice.

---

```
In [ ]:
```

```
In [17]: model = models.vgg13(pretrained=True)
```

```
/home/harkhymadhe/miniforge3/lib/python3.11/site-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.  
  warnings.warn(  
/home/harkhymadhe/miniforge3/lib/python3.11/site-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing `weights=VGG13_Weights.IMAGENET1K_V1`. You can also use `weights=VGG13_Weights.DEFAULT` to get the most up-to-date weights.  
  warnings.warn(msg)
```

```
In [18]: model
```

```
Out[18]: VGG(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace=True)
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU(inplace=True)
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): ReLU(inplace=True)
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (8): ReLU(inplace=True)
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace=True)
    (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (13): ReLU(inplace=True)
    (14): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (15): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (16): ReLU(inplace=True)
    (17): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (18): ReLU(inplace=True)
    (19): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (20): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (21): ReLU(inplace=True)
    (22): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (23): ReLU(inplace=True)
    (24): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(7, 7))
  (classifier): Sequential(
    (0): Linear(in_features=25088, out_features=4096, bias=True)
    (1): ReLU(inplace=True)
    (2): Dropout(p=0.5, inplace=False)
    (3): Linear(in_features=4096, out_features=4096, bias=True)
    (4): ReLU(inplace=True)
    (5): Dropout(p=0.5, inplace=False)
    (6): Linear(in_features=4096, out_features=1000, bias=True)
  )
)
```

```
In [ ]:
```

```
In [19]: class ParticleModel(nn.Module):
    def __init__(self, model, freeze = False, out_features = 2, channels = 2, height = 2, width = 2):
        super().__init__()
        self.backbone = model
        self.freeze = freeze

        self.channels = channels
        self.height = height
        self.width = width

        self.out_features = out_features
        self.layer_norm = nn.LayerNorm([self.channels, self.height, self.width])

        if self.freeze:
            for param in self.backbone.parameters():
                param.requires_grad_(False)

        in_ = self.backbone.classifier[-1].in_features
        self.backbone.classifier[-1] = nn.Linear(in_features = in_, out_features = out_features)

    def forward(self, x):
        x = self.layer_norm(x)
        return self.backbone(x)
```

```
In [20]: def initialize_weights(model):
    for (name, weights) in filter(lambda x: x[1].requires_grad, model.named_parameters()):
        if name.split(".")[1] not in ["fc", "conv1"]:
            continue
        try:
            nn.init.kaiming_normal_(weights)
        except:
            nn.init.normal_(weights, 0., 0.05)

    return model
```

In this notebook, the pretrained weights will be finetuned. This is in contrast to the previous one, where the weights were kept frozen. Also, the learning rate is increased from 1e-4 to 1e-3.

**Update 1:** A learning rate of 1e-3 may be too small for non-frozen weights. I will now attempt to freeze the weights and leave the learning rate as is. Freezing the weights might even speed up training.

**Update 2:** Freezing all the weights seem to have reduced performance. This might be due to the fact that the data we have here is not actually a set of images, even though they seem so. It appears I might have to unfreeze the weights and increase the learning rate a bit.

**Update 3:** Applying the ideas from **Update 2** led to even worse performance! Returning the state of training to **Update 1**...

```
In [21]: EPOCHS = 20
        l2_lambda = 1e-4

        criterion = nn.CrossEntropyLoss().to(DEVICE)

        # Optimizer hyperparameters
        LR = 1e-3
        FACTOR = 100
        AMSGRAD = False
        BETAS = (.9, .999)
        FREEZE = False
```

```
In [22]: model = ParticleModel(
        model = model,
        freeze = FREEZE,
        channels = 3,
        height = 125,
        width = 125
    ).to(DEVICE)
```

```
In [23]: model = initialize_weights(model)
```

```
In [24]: # Instantiate optimizer
        opt = optim.AdamW(
            params = [{
                "params" : model.backbone.parameters(),
                "lr": LR
            }],
            lr=LR/FACTOR,
            amsgrad = AMSGRAD,
            betas = BETAS,
            weight_decay = l2_lambda,
            fused = True
        )

        # scheduler = optim.lr_scheduler.Cos
```

```
In [25]: from sklearn.metrics import accuracy_score
```

```
In [26]: def training_loop(epochs, model, optimizer):
        TRAIN_LOSSES, TEST_LOSSES = [], []
        TRAIN_ACCS, TEST_ACCS = [], []

        for epoch in range(1, epochs + 1):
            train_losses, test_losses = [], []
            train_accs, test_accs = [], []

            model.train() # Set up training mode

            for batch in iter(train_dl):
                # X, y = collate_function(batch)
                X, y = batch
                X, y = X.squeeze().to(DEVICE), y.view(-1).to(DEVICE)
```

```

y_pred = model(X)

train_loss = criterion(y_pred, y.to(torch.long)) # Compare actual
train_loss.backward() # Backpropagate the loss

optimizer.step()
optimizer.zero_grad()

train_losses.append(train_loss.detach().item())

train_acc = accuracy_score(y.cpu().numpy(), y_pred.max(dim = -1))
train_accs.append(train_acc)

# Persist model architecture
torch.save(model.state_dict(), f"epoch_{epoch}_vgg_model.pt")

with torch.no_grad(): # Turn off computational graph
    model.eval() # Set model to evaluation mode
    for batch in iter(test_dl):
        # X_, y_ = collate_function(batch)
        X_, y_ = batch
        X_, y_ = X_.squeeze().to(DEVICE), y_.view(-1).to(DEVICE)

        y_pred_ = model(X_)

        test_loss = criterion(y_pred_, y_.to(torch.long)) # Compare
        test_losses.append(test_loss.item())

        test_acc = accuracy_score(y_.cpu().numpy(), y_pred_.max(dim
        test_accs.append(test_acc)

    avg_train_loss = sum(train_losses) / len(train_losses)
    avg_test_loss = sum(test_losses) / len(test_losses)

    avg_train_acc = sum(train_accs) / len(train_accs)
    avg_test_acc = sum(test_accs) / len(test_accs)

    print(
        f"Epoch: {epoch} | Train loss: {avg_train_loss: .3f} | Test loss
        f"Train accuracy: {avg_train_acc: .3f} | Test accuracy: {avg_test
    )

    TRAIN_LOSSES.append(avg_train_loss)
    TEST_LOSSES.append(avg_test_loss)

    TRAIN_ACCS.append(avg_train_acc)
    TEST_ACCS.append(avg_test_acc)

# Clear CUDA cache
torch.cuda.empty_cache()
torch.clear_autocast_cache()

return {
    "loss": [TRAIN_LOSSES, TEST_LOSSES],
    "accuracy": [TRAIN_ACCS, TEST_ACCS],

```



```
    "model": model
}
```

```
In [ ]: # Train VGG-13 with finetuning
model_results = training_loop(epochs = EPOCHS, optimizer = opt, model = model)
```

```
Epoch: 1 | Train loss: 0.605 | Test loss: 0.560 | Train accuracy: 0.699 |
Test accuracy: 0.719 |
Epoch: 2 | Train loss: 0.573 | Test loss: 0.558 | Train accuracy: 0.715 |
Test accuracy: 0.724 |
Epoch: 3 | Train loss: 0.578 | Test loss: 0.558 | Train accuracy: 0.716 |
Test accuracy: 0.725 |
Epoch: 4 | Train loss: 0.569 | Test loss: 0.557 | Train accuracy: 0.720 |
Test accuracy: 0.727 |
Epoch: 5 | Train loss: 0.563 | Test loss: 0.558 | Train accuracy: 0.725 |
Test accuracy: 0.727 |
Epoch: 6 | Train loss: 0.564 | Test loss: 0.556 | Train accuracy: 0.723 |
Test accuracy: 0.724 |
Epoch: 7 | Train loss: 0.559 | Test loss: 0.558 | Train accuracy: 0.727 |
Test accuracy: 0.725 |
Epoch: 8 | Train loss: 0.559 | Test loss: 0.552 | Train accuracy: 0.728 |
Test accuracy: 0.728 |
Epoch: 9 | Train loss: 0.556 | Test loss: 0.552 | Train accuracy: 0.729 |
Test accuracy: 0.729 |
Epoch: 10 | Train loss: 0.554 | Test loss: 0.551 | Train accuracy: 0.731
| Test accuracy: 0.728 |
Epoch: 11 | Train loss: 0.554 | Test loss: 0.556 | Train accuracy: 0.732
| Test accuracy: 0.728 |
Epoch: 12 | Train loss: 0.553 | Test loss: 0.555 | Train accuracy: 0.732
| Test accuracy: 0.733 |
Epoch: 13 | Train loss: 0.556 | Test loss: 0.552 | Train accuracy: 0.730
| Test accuracy: 0.727 |
Epoch: 14 | Train loss: 0.556 | Test loss: 0.569 | Train accuracy: 0.731
| Test accuracy: 0.723 |
Epoch: 15 | Train loss: 0.550 | Test loss: 0.550 | Train accuracy: 0.734
| Test accuracy: 0.730 |
```

```
In [ ]: # Persist model
torch.save(model_results["model"].state_dict(), "final_epoch_vgg_model.pt")
```

```
In [ ]: def visualize_results(history, key = None):
    if key is not None:
        TRAIN_RESULTS, TEST_RESULTS = history[key]

        plt.figure(figsize = (10, 3))

        plt.plot(range(EPOCHS), TRAIN_RESULTS, label = f"Training {key.capitalize()}")
        plt.plot(range(EPOCHS), TEST_RESULTS, label = f"Test {key.capitalize()}")

        plt.xlabel("Epochs")
        plt.ylabel(key.capitalize())

        plt.title(key.capitalize() + " Evolution for Train and Test Splits",
                  color = "red")

        plt.legend()
        plt.show(); plt.close("all")
```

```

else:
    TRAIN_LOSSES, TEST_LOSSES = history["loss"]
    TRAIN_ACCS, TEST_ACCS = history["accuracy"]

    fig, ax = plt.subplots(1, 2, figsize = (15, 4))

    ax[0].plot(range(EPOCHS), TRAIN_LOSSES, label = "Training Loss")
    ax[0].plot(range(EPOCHS), TEST_LOSSES, label = "Test Loss")

    ax[0].set_xlabel("Epochs")
    ax[0].set_ylabel("Loss")

    ax[0].set_title("Loss Evolution for Train and Test Splits", fontsize

    ax[1].plot(range(EPOCHS), TRAIN_ACCS, label = "Training Accuracy")
    ax[1].plot(range(EPOCHS), TEST_ACCS, label = "Test Accuracy")

    ax[1].set_xlabel("Epochs")
    ax[1].set_ylabel("Accuracy")

    ax[1].set_title("Accuracy Evolution for Train and Test Splits", font

    plt.legend()
    plt.show(); plt.close("all")

return

```

In [ ]:

```

In [ ]: # VGG-13 with finetuning
visualize_results(model_results)

```

In [ ]:

In [ ]: